Lecture 5 Quiz

4/6 points (66%)

Quiz, 6 questions

✖ **Try again once you are ready.**

Required to pass: 80% or higher

You can retake this quiz up to 3 times every 8 hours.

| Back to Week 5 |
| Retake |

✔ 1 / 1
points

1.
For which of the following tasks can we expect that the problem of "dimension hopping" will occur (given that the data is input correctly)? Check all that apply.

☑ Determining whether a given image shows a bike or a car. The bike or car might appear anywhere in the image. The input is the whole set of pixels for the image.

Correct
Dimension hopping occurs when one can take the information contained in the dimensions of some input, and move this between dimensions while not changing the target. The canonical example is taking an image of a handwritten digit and translating it within the image. The dimensions that contain "ink" are now different (they have been moved to other dimensions), however the label we assign to the digit has not changed. Note that this is not something that happens consistently across the dataset, that is we may have a dataset containing two handwritten digits where one is a translated version of the other, however this still does not change the corresponding label of the digits.

☐ Estimating the risk that a patient will develop heart disease given their age, weight, blood pressure, and cholesterol level.

Un-selected is correct

☑ Determining whether a wave has high frequency or low frequency. The input is a set of time values along with their corresponding vertical displacements.

Lecture 5 Quiz

Quiz, 6 questions

4/6 points (66%)

Correct

Dimension hopping occurs when one can take the information contained in the dimensions of some input, and move this between dimensions while not changing the target. The canonical example is taking an image of a handwritten digit and translating it within the image. The dimensions that contain "ink" are now different (they have been moved to other dimensions), however the label we assign to the digit has not changed. Note that this is not something that happens consistently across the dataset, that is we may have a dataset containing two handwritten digits where one is a translated version of the other, however this still does not change the corresponding label of the digits.

☐ Estimating the market price of a house given the number of rooms in the house, the number of schools and the average income of the surrounding neighbourhood, and the average sale price of the surrounding houses.

▲

Un-selected is correct

✔ 1 / 1 points

2.

We have a convolutional neural network for images of 5 by 5 pixels. In this network, each hidden unit is connected to a different 4 x 4 region of the input image:

- The first hidden unit, h_1, is connected to the upper left 4x4 portion of the input image (as shown below).

- The second hidden unit, h_2, is connected to the upper right 4x4 portion of the input image (as shown below).

- The third hidden unit, h_3, is connected to the lower left 4x4 portion of the input image (not shown in the diagram).

- The fourth hidden unit, h_4, is connected to the lower right 4x4 portion of the input image (not shown in the diagram).

Because it's a convolutional network, the weights (connection strengths) are the same for all hidden units: the only difference between the hidden units is that each of them connects to a different part of the input image. In the second diagram, we show the array of weights, which are the same for each of the four hidden units.

Lecture 5 Quiz

Quiz, 6 questions

For h_1, weight w_{11} is connected to the top-left pixel, i.e. x_{11}, while for hidden unit h_2, weight w_{11} connects to the pixel that is one to the right of the top left pixel, i.e. x_{12}.

Lecture 5 Quiz

Quiz, 6 questions

4/6 points (66%)

Imagine that for some training case, we have an input image where each of the black pixels in the top diagram has value 1, and each of the white ones has value 0. Notice that the image shows a "3" in pixels.

The network has no biases. The weights of the network are given as follows:

$w_{11} = 1$	$w_{12} = 1$	$w_{13} = 1$	$w_{14} = 0$
$w_{21} = 0$	$w_{22} = 0$	$w_{23} = 1$	$w_{24} = 0$
$w_{31} = 1$	$w_{32} = 1$	$w_{33} = 1$	$w_{34} = 0$
$w_{41} = 0$	$w_{42} = 0$	$w_{43} = 1$	$w_{44} = 0$

The hidden units are *logistic*.

For the training case with that "3" input image, what is the output of each of the four hidden units?

○ $h_1 = 0.982, h_2 = 0.982, h_3 = 0.982, h_4 = 0.982$

○ $h_1 = 0.982, h_2 = 0.881, h_3 = 1.000, h_4 = 0.982$

◉ $h_1 = 0.982, h_2 = 1.000, h_3 = 0.881, h_4 = 0.982$

▲

Correct

To get the answer we need to start by multiplying the weight matrix by the pixels that it is assigned to in order to get the *total input* to the hidden units. For hidden unit h_4 this would mean:

$$z_4 = w_{11}x_{22} + w_{12}x_{23} + w_{13}x_{24} + w_{14}x_{25}$$
$$+ w_{21}x_{32} + w_{22}x_{33} + w_{23}x_{34} + w_{24}x_{35}$$
$$+ w_{31}x_{42} + w_{32}x_{43} + w_{33}x_{44} + w_{34}x_{45}$$
$$+ w_{41}x_{52} + w_{42}x_{53} + w_{43}x_{54} + w_{44}x_{55}$$

$$= (1)(0) + (1)(0) + (1)(1) + (0)(0)$$
$$+ (0)(1) + (0)(1) + (1)(1) + (0)(0)$$
$$+ (1)(0) + (1)(0) + (1)(1) + (0)(0)$$
$$+ (0)(1) + (0)(1) + (1)(1) + (0)(0)$$
$$= 4$$

Now we pass this through a sigmoid to get
$y_4 = \frac{1}{1+\exp(-z_4)} = \sigma(z_4) = \sigma(4) \approx 0.982$

○ $h_1 = 0.881, h_2 = 0.982, h_3 = 0.982, h_4 = 1.000$

✗　　　0 / 1
　　　points

Lecture 5 Quiz

Quiz, 6 questions

3.

Recall that pooling is the process of combining the outputs of several hidden units to create a single hidden unit. This introduces some invariance to local transformations in the input image.

4/6 points (66%)

In this example we are pooling hidden units h_1, h_2, and h_3. Let's denote the output of hidden unit h_i as y_i. The hidden unit that we're creating by pooling is called h_{pool}, with output y_{pool}. **Sum pooling** is defined as $y_{\text{pool}} = y_1 + y_2 + y_3$.

This form of pooling can be equivalently represented by making h_{pool} a regular hidden unit that takes the output of the other three hidden units, multiplies them by some weights w_1, w_2 and w_3, adds up the resulting numbers and then outputs some function of this sum. This is illustrated in the figure below.

For this to be equivalent to **sum pooling**, what type of neuron should h_{pool} be, and what should the weights be?

○　$w_1 = \frac{1}{3}$, $w_2 = \frac{1}{3}$, $w_3 = \frac{1}{3}$, and h_{pool} is a logistic neuron.

Lecture 5 Quiz

$w_1 = 1, w_2 = 1, w_3 = 1$, and is a logistic neuron.

4/6 points (66%)

Quiz, 6 questions

This should not be selected

○ , , , and is a linear neuron.

○ , , , and is a linear neuron.


✔ 1 / 1
points

4.
Suppose that we have a vocabulary of 3 words, "a", "b", and "c", and we want to predict the next word in a sentence given the previous two words. For this network, we don't want to use feature vectors for words: we simply use the local encoding, i.e. a 3-component vector with one entry being and the other two entries being .

In the language models that we have seen so far, each of the context words has its own dedicated section of the network, so we would encode this problem with two 3-dimensional inputs. That makes for a total of 6 dimensions. For example, if the two preceding words (the "context" words) are "c" and "b", then the input would be . Clearly, the more context words we want to include, the more input units our network must have. More inputs means more parameters, and thus increases the risk of overfitting. Here is a proposal to reduce the number of parameters in the model:

Consider a single neuron that is connected to this input, and call the weights that connect the input to this neuron , and . connects the neuron to the first input unit, connects it to the second input unit, etc. Notice how for every neuron, we need as many weights as there are input dimensions (6 in our case), which will be the number of words times the length of the context. A way to reduce the number of parameters is to *tie* certain weights together, so that they share a parameter. One possibility is to tie the weights coming from input units that correspond to the same word but at different context positions. In our example that would mean that , , and
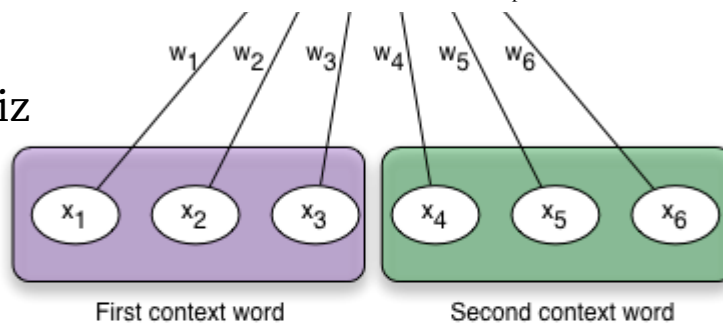
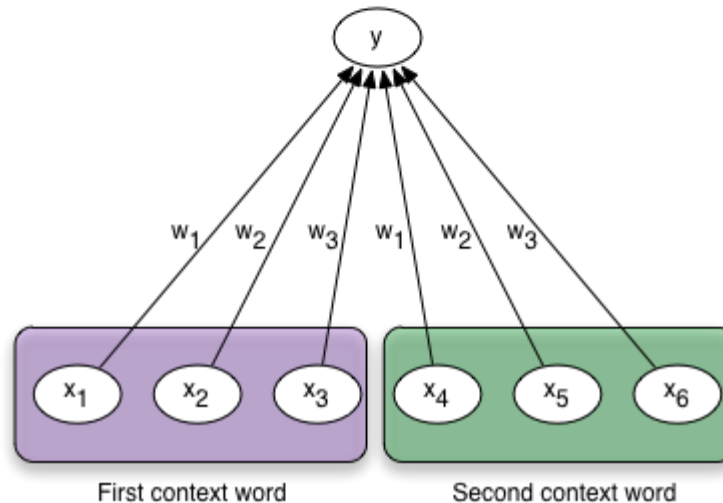(see the "after" diagram below).

Before weight tying

Lecture 5 Quiz

Quiz, 6 questions

After weight tying

Are there any significant problems with this approach?

○ No: the new model after weight tying is an example of a convolutional neural network, and these are more powerful than a non-convolutional network because they are invariant to small transformations in the data.

○ No: this method is an appropriate solution in that it will reduce the number of parameters and therefore always improve generalization.

◉ Yes: the network loses the knowledge of the location at which a context word occurs, and that is valuable knowledge.

Correct

Hint: an equivalent network to the one with weight-tying is one with only 3 inputs and weights (instead of 6 inputs), but instead of storing binary numbers, each dimension now stores the *counts* of each word. That is, each input stores the number of times the word appeared in the sequence. So the input [1 0 0 1 0 0] would be represented as [2 0 0]. What can we say about this network now?

○ Yes: weight tying only makes sense when we are working with images.

Lecture 5 Quiz ✖ •   0 / 1
   points

4/6 points (66%)

Quiz, 6 questions

5.
Let's look at what weight tying does to gradients, computed using the backpropagation algorithm. For this question, our network has three input units,

, two *logistic* hidden units , four input to hidden weights , and two hidden to output weights

. The output neuron is a *linear neuron*, and we are using the *squared error* cost function.

Consider a single training case with target output . The sequence of computations required to compute the error (this is called forward propagation) are as follows:

is the error. Reminder:

and are called the *total inputs* to hidden units and respectively.

Suppose we now decide to tie the weights so that . What is the derivative of the error with respect to ?

Hint: forget about weight tying for a moment and simply compute the derivatives and . Now set .

○

⦿

  ▲

   This should not be selected

○

Lecture 5 Quiz

**4/6 points (66%)**

Quiz, 6 questions

✓   1 / 1
         points

6.

Oh no! Brian's done it again! Claire had a dataset of 28 x 28 pixel handwritten digits nicely prepared to train a neural network, but Brian has gone and accidentally scrambled the images by re-ordering the pixels in some totally meaningless way, and now they can't get the original dataset back! Below is an image of a handwritten '2' before and after being scrambled by Brian.

BeforeAfter

Lecture 5 Quiz

4/6 points (66%)

Luckily, all of the images (in both the training set and the test set) were changed in the same way. For example, if pixels number 1 and number 3 switched places in one image, then they switched places in every other image as well. Because of that, Claire thinks that perhaps she can train a network after all.

Whether Claire is right or not depends largely on the type of neural network that she has in mind. Which of the following neural networks will be at a **disadvantage** because of Brian's mistake? Check all that apply.

- ☑ A convolutional neural network where the size of each weight filter is 10 x 10.

Correct
Imagine that the most important feature for being able to classify this '2' correctly is the little loopy part. A convolutional neural network will normally be able to identify this (provided the weight filter was big enough) because the important piece is located in a small spatial area. Once the dataset has been scrambled, the pixels for the loopy part will be potentially scattered all over the image, so a convolutional network that only has small weight filters will not be able to see enough of the loopy part to identify it.

- ☐ A feed-forward neural network with no hidden layer and linear units (and no convolution).

Un-selected is correct

- ☐ A feed-forward neural network with one hidden layer of logistic units (and no convolution).

Un-selected is correct

- ☑ A convolutional neural network where the size of each weight filter is 8 x 8.

Correct
Imagine that the most important feature for being able to classify this '2' correctly is the little loopy part. A convolutional neural network will normally be able to identify this (provided the weight filter was big enough) because the important piece is located in a small spatial area. Once the dataset has been scrambled, the

Lecture 5 Quiz

Quiz, 6 questions

pixels for the loopy part will be potentially scattered all over the image, so a convolutional network that only has small weight filters will not be able to see enough of the loopy part to identify it.

4/6 points (66%)